AI CAN DO IT · TENCENT CLOUD HACKATHON 2026 · ASPIRE FINTECH TRACK · CASE STUDY: THE I A BRAIN



YieldGPT

One question across Confluence, Jira, Slack and Google Drive. Answered only from what you are allowed to see, and every answer is on the record.

Mohnish · permissions and connectors **Imaad** · retrieval, grounding, UI **Punya** · audit, reliability, deployment

THE PROBLEM

Company knowledge lives on four platforms, each with its own permission rules

C /

Confluence

- Space permissions
- Page restrictions
- Inherited from parent pages

J /

Jira

- Project roles
- Issue security levels
- Inside one project

/

Slack

- Public and private channels
- Direct messages
- Several workspaces

△ /

Drive

- Shared and personal drives
- Folder inheritance
- Link and external shares

> An assistant that copies everything into one index can show people what they were never allowed to open, and nobody can prove afterwards who saw what.

WHAT YIELDGPT DOES

Ask once. Get a cited answer built only from sources you may open right now.

01

Plain-language question

Across all four platforms at once, routed by a planner to the right sources and time window.

02

Each platform decides

Every source is checked live with Confluence, Jira, Slack or Drive itself, as that person. No copied access lists.

03

Everything on the record

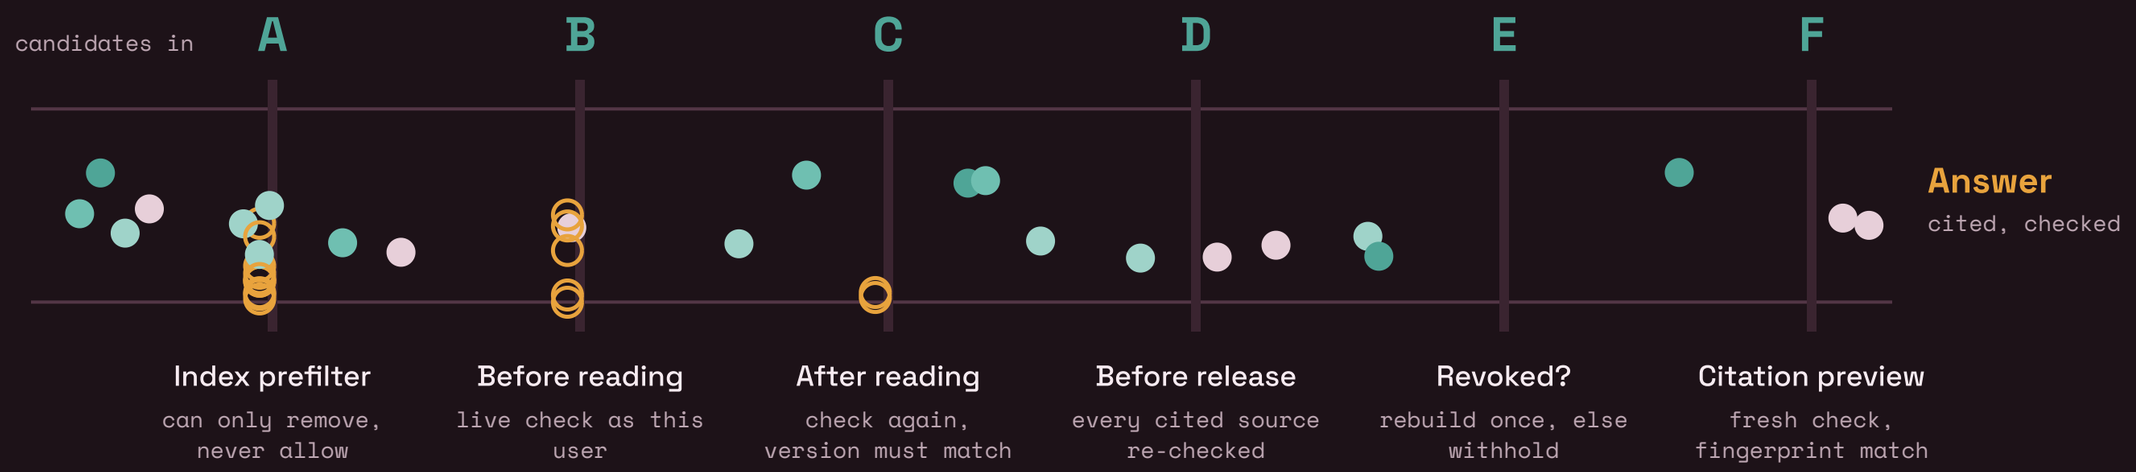
The question, every allow or deny, what the model saw and what was released: hash-chained and checkpointed.

LIVE EXAMPLE

Jira issue PAY-4 sits under an issue security level. Ravi can browse the PAY project and is still denied; Asha, in the security group, gets it. Jira made that call, not us.

HOW IT WORKS

One question, six permission checks

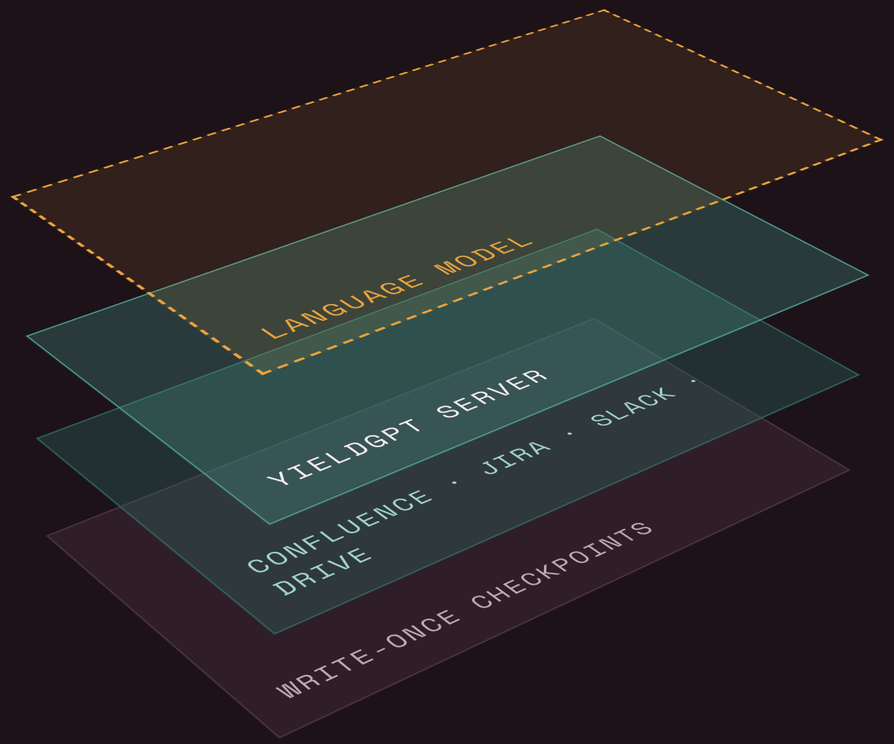


● allowed, still in play ○ denied: dropped, recorded, never shown

One audit row per decision. A denied document looks exactly like one that never existed.

ARCHITECTURE

Four trust boundaries, one direction of trust



UNTRUSTED

The language model. Its output is a proposal, never a decision.

TRUSTED

The YieldGPT server runs every check and writes the audit.

SOURCE OF TRUTH

The platforms themselves decide every permission.

INDEPENDENT

Write-once checkpoint stores the app cannot rewrite.

THE MODEL

The language model is treated as untrusted.

01 It only ever sees passages that already passed checks A to C.

02 It holds no credentials and calls no tools.

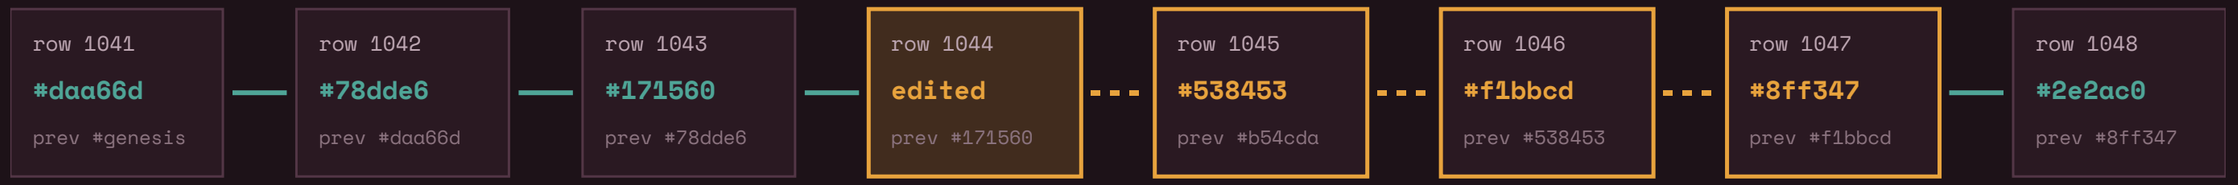
03 Unsupported claims are dropped; citations may only point at sources it was given.

04 Injected instructions in a question are blocked; poisoned passages are quarantined.

In this submission answers are extractive: assembled on the server from authorized passages, word for word, each line cited. The model adapter is optional.

AUDIT AND COMPLIANCE

Every answer is on the record before it leaves



verify: hash mismatch from row 1044, chain broken

Hash-chained log

Each row hashes the one before. Edits, deletions, truncation and a full recompute are all detected. Question and answer text is encrypted at rest.

No checkpoint, no answer

The answer is held until its audit rows are checkpointed. If checkpointing fails, the user gets a refusal, never an unrecorded answer.

Two write-once anchors

S3 Object Lock in compliance mode, written to two stores. A forged anchor disagrees with the other and verification fails.

PROOF · LIVE

Verified on a real Atlassian Cloud site

84/84

permission decisions matched Atlassian: every persona against every seeded page and issue

14/14

brief scenarios passed as each persona: unified query, negative cases, audit inquiry

21/21

lifecycle steps through the real pipeline: create, edit, revoke, delete, inherited restriction, status and comment

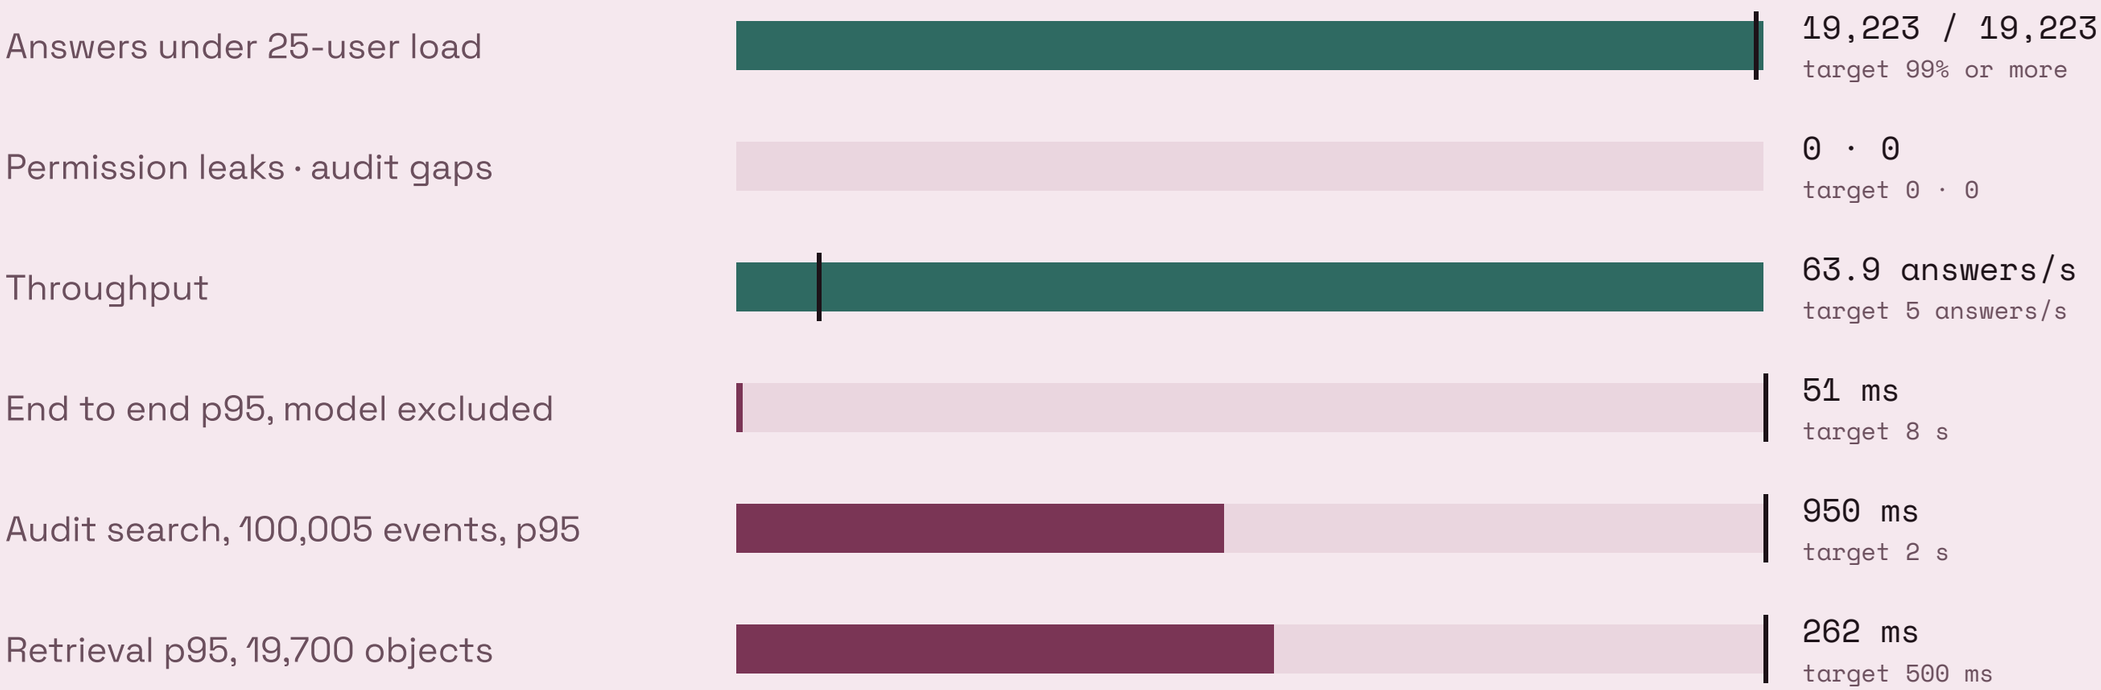
Revoked user denied on the **very next query**, no reindex: 0.27 s Confluence, 0.18 s Jira. Denial held 3.2 s / 3.0 s after Atlassian confirmed, never re-allowed.

A denial looks identical to a document that never existed. A child page inherits its parent's restriction. Deleted objects are not served on the next query. Issue security levels enforced by Jira itself.

Live runs, 13 to 14 Oct 2026 · Atlassian Cloud Standard trial · four teammates as personas · single dated runs, not throughput · Slack and Drive simulated in these runs

PROOF · UNDER LOAD

25 users for 5 minutes: zero leaks, zero audit gaps



10 to 12 Oct 2026 · Apple Silicon laptops · fixture data and stub model; search and retrieval synthetic · end to end includes write-once checkpoints (MinIO) · real-model time not included · bars against target

PROOF · ANSWER QUALITY

Same questions, original pipeline vs YieldGPT

Questions passing	Required-source recall	Grounding cases	False accepts
~~50/65~~	~~85.7%~~	~~15/28~~	~~13~~
65/65	**100%**	**28/28**	**0**

Neither pipeline exposed an unauthorized citation. 440+ automated tests run on every pull request in CI, in two connector modes.

Fixture data · offline comparison on identical cases · human scoring of the 160-question bank is still in progress

TENCENT ECOSYSTEM

YieldGPT inside the tools Tencent teams already use

WORKING

CodeBuddy and WorkBuddy

An MCP connector lets the CodeBuddy agent call YieldGPT and open citations. Identity comes from config, so every check still applies.

ADAPTER READY

Tencent Hunyuan

The adapter speaks the OpenAI-compatible API Hunyuan offers and is tested end to end. This submission answers in extractive mode; switching Hunyuan on is configuration.

IN PROGRESS

Tencent Cloud

Docker deployment with secrets from the environment only, and a public demo URL. Tencent COS is a natural second checkpoint anchor.

HONEST STATUS

What is live, and what is still pending

LIVE AND VERIFIED

Confluence and Jira on a real Atlassian site

84/84 decisions · 14/14 scenarios · 21/21 lifecycle

Write-once checkpoints on two stores

S3 Object Lock, compliance mode

YieldGPT inside CodeBuddy

Full test suite in CI on every pull request

PENDING

Hunyuan switched on and evaluated

answers are extractive by design in this submission

Live Slack workspaces and live Google Drive

Public URL on Tencent Cloud

Two-reviewer human scoring

Every number in this deck carries a date and a label: live, fixture or synthetic.

YIELDGPT



Ask once.
See only what you may.
Prove it afterwards.

github.com/Imaad001M/Tencent---YieldGPT

Mohnish · Imaad · Punya